

Fall 21. Nov 01, 2021. 12:18 pm EST. This notebook integrates two - the first part introduces basics of Pandas; the lower part suggests applying pandas/python as first steps getting to know our data.  
See also <https://pandas.pydata.org>

This week, we get ready for Project 2, review import first steps for analysis and stats with Pandas, optionally look at large data architectures, and underscore how to think about integrating pandas & python in approaching unknown data sets.



---

### This week's topics:

This week we look at Pandas. It is important, too, to begin to contextualize these skills in analysis, practice, associate programming with research practices, and more, as rationales for many of Panda's tools.

---

## Agenda

### 1. Data Exploration & Analysis - in General

2. Overview: Integrating Python and Pandas into this work
  3. Exploration: transforming data
  4. Optional: architectures and work practices; data wrangling review pages
  5. Pandas: series, data frames, panel
  6. Pandas Code Samples
  7. Optional Pandas-on-the-job with visuals
  8. Optional Review
  9. Breakout Room
  10. Project 2
-

# 1. Data Exploration & Analysis

## Data Exploration:

- For data integrity
- To develop questions based on the variables
- To break your model - better now than in production

## Data Analysis:

- Answer a research question or hypothesis
- Usually involves complex math, modeling statistics
- Likely to combine datasets
- Explore data by collapsing in groups in various ways
- Some functions are useful in exploration & analysis

---


### Why use Pandas?

Pandas and other tools help in machine learning, scaling data, multiple regression, and more! Usually work together with numpy, scipy, matplotlib, scikit-learn, pysqlite3, psycopg2, and others. These two sites introduce a *lot* of features for python to read lots of data sources and more efficient problem-solving techniques: [https://pandas.pydata.org/pandas-docs/stable/getting\\_started/10min.html](https://pandas.pydata.org/pandas-docs/stable/getting_started/10min.html) ([https://pandas.pydata.org/pandas-docs/stable/getting\\_started/10min.html](https://pandas.pydata.org/pandas-docs/stable/getting_started/10min.html)) and [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/io.html#io-hdf5](https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html#io-hdf5) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/io.html#io-hdf5](https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html#io-hdf5)).

---

## 2. Overview Exploration & Analysis: Discuss

Take time to get to know your data - the domain & range, parametric or non-parametric; explore the measures of central tendency ... Here are some handy commands:



- `value_counts()`
- `describe()`
- `min(), max(), isnull()`
- Plot your data during exploration, too, not just during analysis.
- Consider the source(s) of your data and research the topic:
  - basic research methods require looking at threats to validity,
  - cross-validating the data,
  - issues of research "bias",
  - lack of precision in definitions (e.g., mismatched metadata when combining data)
  - look for any professional/industrial gold standards for measurement
- what might be confounding events in your data?

You might want to learn more about the expectations of "research" by reading Booth, et al., *Craft of research*.

---

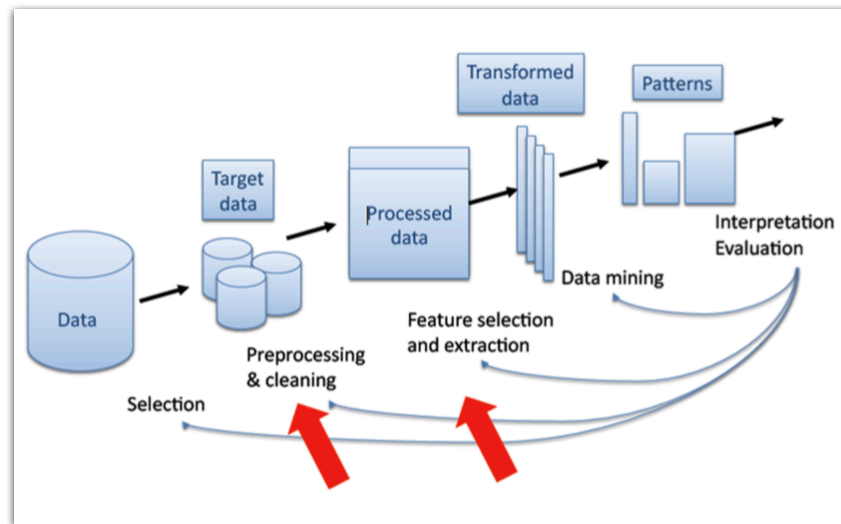
## 3. Exploration & Analysis: Transformations for analysis

- **filter** based on some condition (e.g., too high values? duplicated data?)
  - **create new** columns, when necessary
  - **aggregate or collapse** by groups
  - **join** two+ datasets together
-

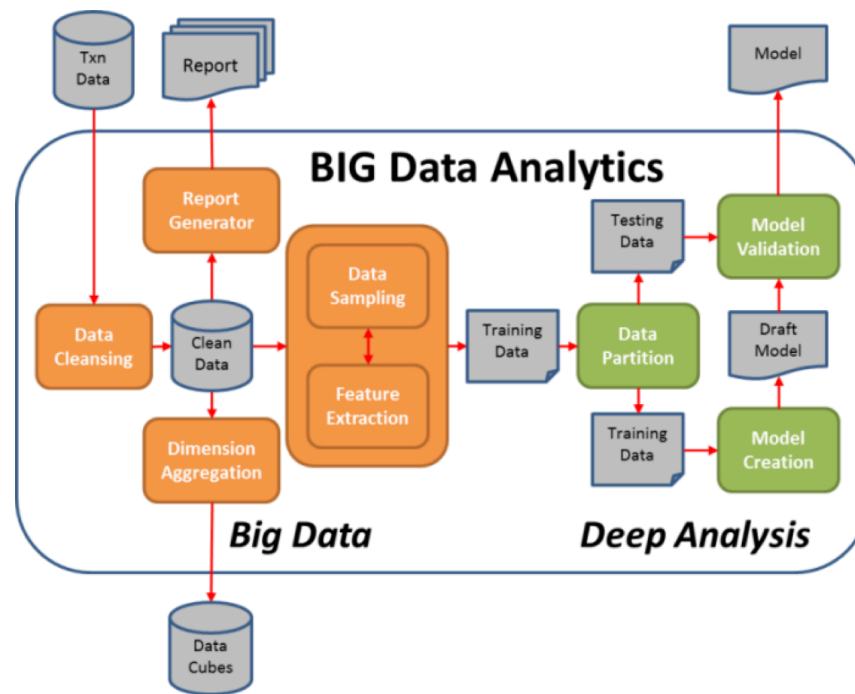
## 4. Optional Discussion

Applying exploration & analysis in larger settings: data mining, big data analytics, Hadoop/activities and helpers.

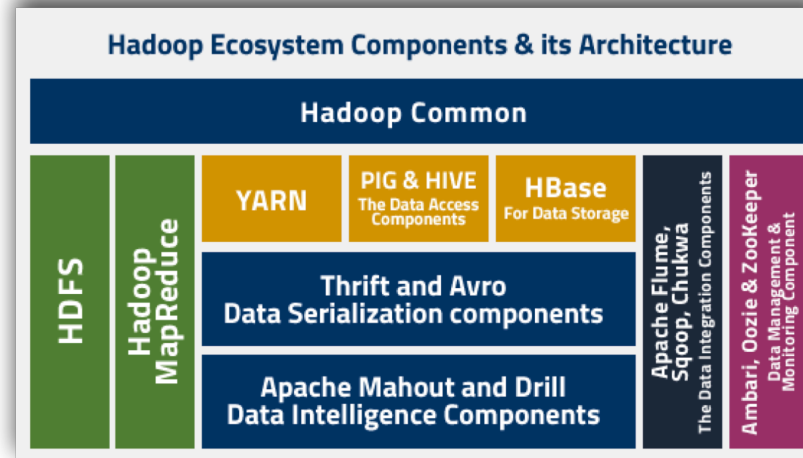
Notice that as datasets get larger and exploration/analysis more complex, we identify specific work behaviors, warehouses, and architectures/other programming languages as part of the job. See the Big Data book in the optional resources.



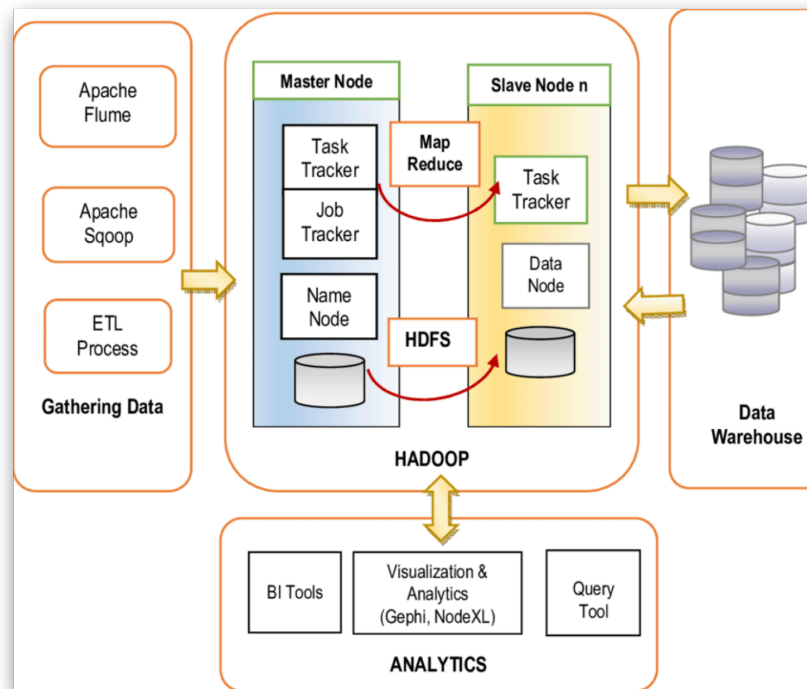
Data Mining Activities



Big Data



Hadoop

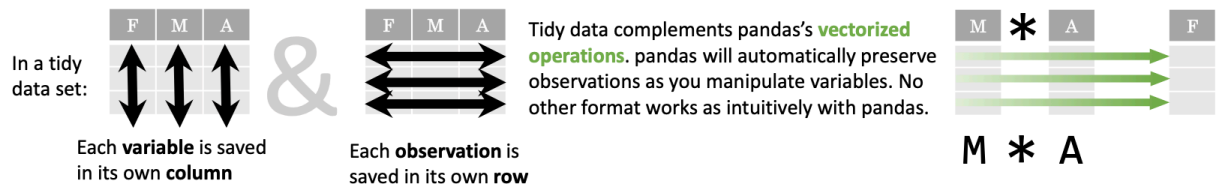


Tools

## Data Wrangling with pandas Cheat Sheet

<http://pandas.pydata.org>

### Tidy Data – A foundation for wrangling in pandas



### Syntax – Creating DataFrames

```
df = pd.DataFrame(
    {"a": [4, 5, 6],
     "b": [7, 8, 9],
     "c": [10, 11, 12]},
    index = [1, 2, 3])
Specify values for each column.

df = pd.DataFrame(
    [[4, 7, 10],
     [5, 8, 11],
     [6, 9, 12]])
```

### Reshaping Data – Change the layout of a data set

```
pd.melt(df)
Gather columns into rows.

df.pivot(columns='var', values='val')
Spread rows into columns.
```

```
df.sort_values('mpg')
Order rows by values of a column (low to high).

df.sort_values('mpg', ascending=False)
Order rows by values of a column (high to low).

df.rename(columns = {'y': 'year'})
Rename the columns of a DataFrame

df.sort_index()
Sort the index of a DataFrame

df.reset_index()
Reset index of DataFrame to row numbers, moving
index to columns.
```

```
index=[1, 2, 3],
columns=['a', 'b', 'c'])
```

Specify values for each row.

	a	b	c
n			
d	1	4	7
e	2	6	9

```
df = pd.DataFrame(
    {"a": [4, 5, 6],
     "b": [7, 8, 9],
     "c": [10, 11, 12]},
    index = pd.MultiIndex.from_tuples(
        [('d', 1), ('d', 2), ('e', 2)],
        names=['n', 'v']))
```

Create DataFrame with a MultiIndex

## Method Chaining

Most pandas methods return a DataFrame so that another pandas method can be applied to the result. This improves readability of code.


```
df = (pd.melt(df)
      .rename(columns={
          'variable': 'var',
          'value': 'val'})
      .query('val >= 200'))
```

`pd.concat([df1, df2])`  
Append rows of DataFrames

`pd.concat([df1, df2], axis=1)`  
Append columns of DataFrames

`df.drop(columns=['length', 'height'])`  
Drop columns from DataFrame

## Subset Observations (Rows)



`df[df.Length > 7]`  
Extract rows that meet logical criteria.

`df.drop_duplicates()`  
Remove duplicate rows (only considers columns).

`df.head(n)`  
Select first n rows.

`df.tail(n)`  
Select last n rows.

`df.sample(frac=0.5)`  
Randomly select fraction of rows.

`df.sample(n=10)`  
Randomly select n rows.

`df.iloc[10:20]`  
Select rows by position.

`df.nlargest(n, 'value')`  
Select and order top n entries.

`df.nsmallest(n, 'value')`  
Select and order bottom n entries.

## Subset Variables (Columns)



`df[['width', 'length', 'species']]`  
Select multiple columns with specific names.

`df['width']` or `df.width`  
Select single column with specific name.

`df.filter(regex='regex')`  
Select columns whose name matches regular expression `regex`.

regex (Regular Expressions) Examples	
<code>'\.'</code>	Matches strings containing a period '.'
<code>'Length\$'</code>	Matches strings ending with word 'Length'
<code>'^Sepal'</code>	Matches strings beginning with the word 'Sepal'
<code>'^x[1-5]\$'</code>	Matches strings beginning with 'x' and ending with 1,2,3,4,5
<code>'^(?!Species\$).*'</code>	Matches strings except the string 'Species'

`df.loc[:, 'x2': 'x4']`  
Select all columns between x2 and x4 (inclusive).

`df.iloc[:, [1, 2, 5]]`  
Select columns in positions 1, 2 and 5 (first column is 0).

`df.loc[df['a'] > 10, ['a', 'c']]`  
Select rows meeting logical condition, and only the specific columns.

Logic in Python (and pandas)		
<code>&lt;</code>	Less than	<code>!=</code> Not equal to
<code>&gt;</code>	Greater than	<code>df.column.isin(values)</code> Group membership
<code>==</code>	Equals	<code>pd.isnull(obj)</code> Is NaN
<code>&lt;=</code>	Less than or equals	<code>pd.notnull(obj)</code> Is not NaN
<code>&gt;=</code>	Greater than or equals	<code>&amp;,  , ~, ^, df.any(), df.all()</code> Logical and, or, not, xor, any, all

<http://pandas.pydata.org/> This cheat sheet inspired by Rstudio Data Wrangling Cheatsheet (<https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>) Written by Irv Lustig, Princeton Consultants



## Summarize Data

**df['w'].value\_counts()**  
Count number of rows with each unique value of variable

**len(df)**  
# of rows in DataFrame.

**df['w'].nunique()**  
# of distinct values in a column.


**df.describe()**  
Basic descriptive statistics for each column (or GroupBy)



pandas provides a large set of **summary functions** that operate on different kinds of pandas objects (DataFrame columns, Series, GroupBy, Expanding and Rolling (see below)) and produce single values for each of the groups. When applied to a DataFrame, the result is returned as a pandas Series for each column. Examples:

<b>sum()</b> Sum values of each object.	<b>min()</b> Minimum value in each object.
<b>count()</b> Count non-NA/null values of each object.	<b>max()</b> Maximum value in each object.
<b>median()</b> Median value of each object.	<b>mean()</b> Mean value of each object.
<b>quantile([0.25,0.75])</b> Quantiles of each object.	<b>var()</b> Variance of each object.
<b>apply(function)</b> Apply function to each object.	<b>std()</b> Standard deviation of each object.

## Group Data



**df.groupby(by="col")**  
Return a GroupBy object, grouped by values in column named "col".

**df.groupby(level="ind")**  
Return a GroupBy object, grouped by values in index level named "ind".

All of the summary functions listed above can be applied to a group. Additional GroupBy functions:

**size()**  
Size of each group.

**agg(function)**  
Aggregate group using function.

## Windows

**df.expanding()**  
Return an Expanding object allowing summary functions to be applied cumulatively.

**df.rolling(n)**  
Return a Rolling object allowing summary functions to be applied to windows of length n.

## Handling Missing Data

**df.dropna()**  
Drop rows with any column having NA/null data.

**df.fillna(value)**  
Replace all NA/null data with value.

## Make New Columns



**df.assign(Area=lambda df: df.Length\*df.Height)**  
Compute and append one or more new columns.

**df['Volume'] = df.Length\*df.Height\*df.Depth**  
Add single column.

**pd.qcut(df.col, n, labels=False)**  
Bin column into n buckets.



pandas provides a large set of **vector functions** that operate on all columns of a DataFrame or a single selected column (a pandas Series). These functions produce vectors of values for each of the columns, or a single Series for the individual Series. Examples:


<b>max(axis=1)</b> Element-wise max.	<b>min(axis=1)</b> Element-wise min.
<b>clip(lower=-10, upper=10)</b> Trim values at input thresholds	<b>abs()</b> Absolute value.

The examples below can also be applied to groups. In this case, the function is applied on a per-group basis, and the returned vectors are of the length of the original DataFrame.


<b>shift(1)</b> Copy with values shifted by 1.	<b>shift(-1)</b> Copy with values lagged by 1.
<b>rank(method='dense')</b> Ranks with no gaps.	<b>cumsum()</b> Cumulative sum.
<b>rank(method='min')</b> Ranks. Ties get min rank.	<b>cummax()</b> Cumulative max.
<b>rank(pct=True)</b> Ranks rescaled to interval [0, 1].	<b>cummin()</b> Cumulative min.
<b>rank(method='first')</b> Ranks. Ties go to first value.	<b>cumprod()</b> Cumulative product.

## Plotting

**df.plot.hist()**  
Histogram for each column



**df.plot.scatter(x='w', y='h')**  
Scatter chart using pairs of points



## Combine Data Sets

adf		bdf	
x1	x2	x1	x3
A	1	A	T
B	2	B	F
C	3	D	T

+

=

### Standard Joins

**pd.merge(adf, bdf, how='left', on='x1')**  
Join matching rows from bdf to adf.

**pd.merge(adf, bdf, how='right', on='x1')**  
Join matching rows from adf to bdf.

**pd.merge(adf, bdf, how='inner', on='x1')**  
Join data. Retain only rows in both sets.

**pd.merge(adf, bdf, how='outer', on='x1')**  
Join data. Retain all values, all rows.

### Filtering Joins

**adf[adf.x1.isin(bdf.x1)]**  
All rows in adf that have a match in bdf.

**adf[~adf.x1.isin(bdf.x1)]**  
All rows in adf that do not have a match in bdf.

ydf		zdf	
x1	x2	x1	x2
A	1	B	2
B	2	C	3
C	3	D	4

+

=

### Set-like Operations

**pd.merge(ydf, zdf)**  
Rows that appear in both ydf and zdf (Intersection).

**pd.merge(ydf, zdf, how='outer')**  
Rows that appear in either or both ydf and zdf (Union).

**pd.merge(ydf, zdf, how='outer', indicator=True)**  
**.query('\_merge == "left\_only"')**  
**.drop(columns=['\_merge'])**  
Rows that appear in ydf but not zdf (Setdiff).

End of the optional section.

---

---

## 5. Exploration & Analysis: Transformations for analysis

Type	Description	Example
Series	1D labeled homogeneous array, size is immutable.	<pre>import pandas as pd import numpy as np data = np.array(['a', 'b', 'c']) s = pd.Series(data) print(s)</pre>
Data Frames	General 2d labeled, size-mutable tabular structure with potentially heterogeneous-typed columns	<pre>import pandas as pd data = [1, 2, 3, 4, 5] df = pd.DataFrame(data)</pre>
Panel	General 3d labeled, size-mutable array	<pre>import pandas as pd import numpy as np data = np.random.rand(2, 4, 5) p = pd.Panel(data)</pre>

---

## 6. Examples Demonstrating Pandas

Pandas components are `Series` (a column) and `DataFrame`, a multi-dimensional table made up of `Series`.

Series			Series			DataFrame		
	apples			oranges			apples	oranges
0	3	+	0	0	=	0	3	0
1	2		1	3		1	2	3
2	0		2	7		2	0	7
3	1		3	2		3	1	2

Use pandas to create a frame (that looks like a combo of json and a dictionary):

```
data = {  
    'cats': [1, 4, 2, 8],  
    'dogs': [1, 3, 3, 2]  
}  
pets = pd.DataFrame(data)
```

```
In [22]: 1 import pandas as pd
          2
          3 data = {
          4     'cats': [1, 4, 2, 8],
          5     'dogs': [1, 3, 3, 2]
          6 }
          7 pets = pd.DataFrame(data)
          8 pets
```

Out[22]:

	<b>cats</b>	<b>dogs</b>
<b>0</b>	1	1
<b>1</b>	4	3
<b>2</b>	2	3
<b>3</b>	8	2

```
In [23]: 1 import pandas as pd
          2
          3 data = {
          4     'cats': [123, 14, 42, 8],
          5     'dogs': [81, 33, 113, 132]
          6 }
          7 new_pets = pd.DataFrame(data, index=['Paris', 'Boston', 'Rome', 'Monterey'])
          8 new_pets
```

Out[23]:

	<b>cats</b>	<b>dogs</b>
<b>Paris</b>	123	81
<b>Boston</b>	14	33
<b>Rome</b>	42	113
<b>Monterey</b>	8	132

```
In [24]: 1 """ if the data set were really large
          2         it'd be nice to locate by key
          3         """
          4 new_pets.loc['Paris']
```

```
Out[24]: cats      123
         dogs       81
         Name: Paris, dtype: int64
```

---

## Building up class examples

In these examples, we want to build up from basics to increasingly complex uses of pandas for coding practices.

---

In [25]:

```
1 import pandas as pd
2
3 data = pd.Series([1,2,3,4,6,0,85,45,7,53,321,4,32,2355,6])
4
5 # select values from data; those < 10
6 print("select data < 10: ", data[data < 10])
7 print("range of data < 10 > 5: ", data[(data < 10) & (data > 5)])
8 print("_"*50,"\n")
9 # same result but "chained"
10 print("same results but 'chained': ", data[(data < 10) & (data > 5)])
11 print("same results without using &: ", data[data < 10][data > 5])
```

```
select data < 10:  0      1
1      2
2      3
3      4
4      6
5      0
8      7
11     4
14     6
dtype: int64
range of data < 10 > 5:  0      1
1      2
2      3
3      4
5      0
11     4
dtype: int64
```

---

```
same results but 'chained':  4      6
8      7
14     6
dtype: int64
same results without using &:  4      6
8      7
14     6
dtype: int64
```

---

## Slicing and Manipulating

Here we'll make a copy first (for safekeeping) and explore using head and slicing and replacing, etc.



In [26]:

```
1 d5 = data.head().copy()
2 print("d5: ", d5)
3 d6 = data.head(70).copy() # specify the size of the head
4 print("d6: ", d6)
5
6 print(data[0:10:2]) # slide the data
7
8 data[0] = 10000 # value replacement
```

d5: 0 1

1 2

2 3

3 4

4 6

dtype: int64

d6: 0 1

1 2

2 3

3 4

4 6

5 0

6 85

7 45

8 7

9 53

10 321

11 4

12 32

13 2355

14 6

dtype: int64

0 1

2 3

4 6

6 85

8 7

dtype: int64

---

## Using any and all condition testing

```
In [27]: 1 print("any: ",data[data < 10].any()) # should return true
          2 print("all: ",data[data < 10].all()) # false
          3 print("alternative: ",(data > 10000).any()) # alternative
          4 print("all, diff syntax", (data > 0).all) # false
```

```
any: True
all: False
alternative: False
all, diff syntax <bound method Series.all of 0      True
1      True
2      True
3      True
4      True
5      False
6      True
7      True
8      True
9      True
10     True
11     True
12     True
13     True
14     True
dtype: bool>
```

---

## Get to know your data for basic stats

```
In [28]: 1 print("Length of the data set:", len(data))
          2
          3 # individual measurements - can use the describe() for the whole set
```

```

3 # Individual measurements - can use the describe() for the whole set
4 print("mean: ", data.mean())
5 print("mode: ", data.mode())
6 print("median: ", data.median())
7 print("count: ", data.count())
8 print("std: ", data.std())
9 print("unique: ", data.unique())
10
11 print("\ndescribe: ", data.describe())
12 print("\n\nvalue_counts and shape of the data.\nNote that shape is an attribute, not a method")
13 print("\tValue counts: ", data.value_counts())
14 print("\tShape: ", data.shape)

```

Length of the data set: 15

mean: 861.5333333333333

mode: 0 4

1 6

dtype: int64

median: 7.0

count: 15

std: 2598.470975309097

unique: [10000 2 3 4 6 0 85 45 7 53 321 32  
2355]

describe: count 15.000000

mean 861.533333

std 2598.470975

min 0.000000

25% 4.000000

50% 7.000000

75% 69.000000

max 10000.000000

dtype: float64

value\_counts and shape of the data.

Note that shape is an attribute, not a method

Value counts: 6 2

4 2

85 1

53 1

```
2355      1
10000     1
45         1
7          1
32         1
3          1
2          1
321        1
0          1
dtype: int64
Shape: (15,)
```

---

## Adding, editing data

```
In [29]: 1 print("by single index names: ", data[10]) # by single index names
          2
          3
          4 print("look up by index location, using dict stye[]")
          5 print(data.iloc[ [0,3] ])
```

```
by single index names: 321
look up by index location, using dict stye[]
0      10000
3         4
dtype: int64
```

---

## Filling data cells

It's often useful to fill cells with some placeholder data and interpolate missing values. This is common in k-nearest neighbor and any situation to minimize need to check for data-oriented problems when running your scripts.

```
In [30]: 1 combo = pd.Series([0,0,0,0,0])
2 new_combo = pd.Series([0,0,0,0,0])
3 new_combination = pd.Series([0,0,0,0,0])
4
5 # set the fill value.
6 print("combo after reindex")
7 combo.reindex([0, 2, 15, 21], fill_value = 0)
8 print(combo)
9
10 # fill the NaNs
11 new_combo.fillna(0)
12
13 # forward and back fill to guess at missing values.
14 new_combo.fffll()
15
16 new_combo.bfill()
17
18 new_combination.interpolate() # one of many techniques.
```

```
combo after reindex
```

```
0    0
1    0
2    0
3    0
4    0
dtype: int64
```

```
Out[30]: 0    0
1    0
2    0
3    0
4    0
dtype: int64
```

---

**Integrating data** - input/output to csv, sql, json

Typical to read in data - sometimes with the metadata (such as an SQL table header) or not - change to .json and import to a dictionary.

Examples:

```
myframe = pd.read_csv('mydata.csv', index_col = 0)
myframe = pd.read_json('mydata.json') . May need the orient keyword.
```

If you're reading from sql, you need to **import sqlite3**. Then you can establish a connection to the database and then select data from your table(s), e.g., `my_connection = sqlite3.connect("mydatabase.db")`

**Output/convert to different data types:**

```
dataframe.to_csv('newfile.csv')
dataframe.to_json('newfile.json')
dataframe.to_sql('newfile', my_connection)
```

---

## Get to know your data first

1. What's the dimensions of our data? `my_dataframe.shape`
2. Get some info about your data: `my_dataframe.info()`
3. See the first 5 rows (the head) of your data: `my_dataframe.head()`
4. last rows: `my_dataframe.tail()`
5. remove duplicates `drop_duplicates()`
6. copy the frame to a new one, or overwrite it `my_dataframe.drop_duplicates(inplace = True)` (We can keep the first of the duplicates (the default) or drop all duplicates: `temp_df = my_dataframe.append(my_dataframe)` to make a copy; `temp_df.drop_duplicates(inplace=True, keep=False)` would overwrite the df and remove all duplicates (keep=False).

Movie data from [here \(https://gist.github.com/tiangechen/b68782efa49a16edaf07dc2cdaa855ea#file-movies-csv\)](https://gist.github.com/tiangechen/b68782efa49a16edaf07dc2cdaa855ea#file-movies-csv)  
`movies_df = pd.read_csv("movies.csv", index_col="Film")`

```
In [31]: 1 movies_df = pd.read_csv("files/movies.csv", index_col="Film")
          2 movies_df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 77 entries, Zack and Miri Make a Porno to (500) Days of Summer
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Genre                  77 non-null    object
1   Lead Studio            77 non-null    object
2   Audience score %       77 non-null    int64
3   Profitability          77 non-null    float64
4   Rotten Tomatoes %      77 non-null    int64
5   Worldwide Gross        77 non-null    object
6   Year                   77 non-null    int64
dtypes: float64(1), int64(3), object(3)
memory usage: 4.8+ KB
```

```
In [32]: 1 movies_df = pd.read_csv("files/movies.csv", index_col="Film")
          2 movies_df.shape
```

```
Out[32]: (77, 7)
```

In [33]: 1 movies\_df.head()

Out[33]:

	Genre	Lead Studio	Audience score %	Profitability	Rotten Tomatoes %	Worldwide Gross	Year
Film							
<b>Zack and Miri Make a Porno</b>	Romance	The Weinstein Company	70	1.747542	64	\$41.94	2008
<b>Youth in Revolt</b>	Comedy	The Weinstein Company	52	1.090000	68	\$19.62	2010
<b>You Will Meet a Tall Dark Stranger</b>	Comedy	Independent	35	1.211818	43	\$26.66	2010
<b>When in Rome</b>	Comedy	Disney	44	0.000000	15	\$43.04	2010
<b>What Happens in Vegas</b>	Comedy	Fox	72	6.267647	28	\$219.37	2008

In [34]: 1 movies\_df.tail()

Out[34]:

	Genre	Lead Studio	Audience score %	Profitability	Rotten Tomatoes %	Worldwide Gross	Year
Film							
<b>Across the Universe</b>	romance	Independent	84	0.652603	54	\$29.37	2007
<b>A Serious Man</b>	Drama	Universal	64	4.382857	89	\$30.68	2009
<b>A Dangerous Method</b>	Drama	Independent	89	0.448645	79	\$8.97	2011
<b>27 Dresses</b>	Comedy	Fox	71	5.343622	40	\$160.31	2008
<b>(500) Days of Summer</b>	comedy	Fox	81	8.096000	87	\$60.72	2009



---

## Discussion

How might you *read in* your data and how *why* might you use Pandas and these commands?

---

---

### Optional

## Getting to you know your data - visually


Next class we look at data visualization. This section is an optional teaser integrating the exploration and analysis possible with pandas with another language of science, graphics.

After you've ingested and gotten to know your data, you'll want to explore it - checking the data distribution, domain and range, and test your data against some specific test or idea.

---

### Normal Distribution


```
In [35]: 1 import numpy
          2 import matplotlib.pyplot as plt
          3
          4 x = numpy.random.normal(5.0, 1.0, 1000000)
          5
          6 plt.hist(x, 100)
          7 plt.show()
```



### Scatter Plot

Two arrays with 1000 random numbers; first array has the mean set to 5.0 with std dev of 1.0; 2nd array has mean set to 10 with std dev 2.0

```
In [36]: 1 # see also https://matplotlib.org/stable/users/dflt\_style\_changes.html
2
3 import numpy
4 import matplotlib.pyplot as plt
5
6 x = numpy.random.normal(5.0, 1.0, 1000)
7 y = numpy.random.normal(10.0, 2.0, 1000)
8
9 plt.scatter(x, y)
10 plt.show()
```




## Linear Regression

Here the x-axis represents some value, say age, and y is the height in millimeters.

2nd plot includes the correlation coefficient ( $r$ ).

[See the SciPy.org Tutorial \(https://docs.scipy.org/doc/scipy/reference/tutorial/index.html\)](https://docs.scipy.org/doc/scipy/reference/tutorial/index.html)

```
In [37]: 1 import matplotlib.pyplot as plt
2
3 x = [5,7,8,7,2,17,2,9,4,11,12,9,6]
4 y = [99,86,87,88,111,86,103,87,94,78,77,85,86]
5
6 plt.scatter(x, y)
7 plt.show()
8
9 # can you determine r and r^2?
```




```
In [38]: 1 import matplotlib.pyplot as plt
2 from scipy import stats
3
4 x = [5,7,8,7,2,17,2,9,4,11,12,9,6]
5 y = [99,86,87,88,111,86,103,87,94,78,77,85,86]
6
7 slope, intercept, r, p, std_err = stats.linregress(x, y)
8
9 print(r)
```

-0.758591524376155

In [39]:

```
1 import numpy
2 import matplotlib.pyplot as plt
3
4 x = [1,2,3,5,6,7,8,9,10,12,13,14,15,16,18,19,21,22]
5 y = [100,90,80,60,60,55,60,65,70,70,75,76,78,79,90,99,99,100]
6
7 mymodel = numpy.poly1d(numpy.polyfit(x, y, 3))
8
9 myline = numpy.linspace(1, 22, 100)
10
11 plt.scatter(x, y)
12 plt.plot(myline, mymodel(myline))
13 plt.show()
```



---

Example: In this example, read data from a .csv file into a DataFrame to measure carbon dioxide. The x var is independent; y is dependent.

The human\_subjects.csv file represents a paired-sample groups of people who are overweight, exercised a set number of hours, and to predict their blood oxygen saturation.

```
In [40]: 1 import pandas
          2 from sklearn import linear_model
          3
          4 df = pandas.read_csv("files/human_subjects.csv")
          5
          6 X = df[['hours', 'weight']]
          7 y = df['02']
          8
          9 regr = linear_model.LinearRegression()
         10 regr.fit(X, y)
         11
         12 predicted02 = regr.predict([[25, 190]])
         13
         14 print(predicted02)
```

```
[95.54921072]
```

---

# Review

This section demonstrates a kind of first on-the-job looking at our data.

Pandas Exploration ... and Review  
Say you're working with some data and you want to explore using a Series or a DataFrame...

## Series

```
In [41]: 1 import pandas as pd
          2 s = pd.Series([100,500,400,300,222])
          3 print(s)
          4 # lets see the data as a list by index
          5 print( list(s.index) )
```

```
0    100
1    500
2    400
3    300
4    222
dtype: int64
[0, 1, 2, 3, 4]
```

## Labeling of data

makes it a lot easier to keep track of things.

```
In [42]: 1 s.index = ['Tom','Jane','Ming','Felicia','Toby']
          2 s.name = "Cousins"
          3 print(s)
```

```
Tom          100
Jane          500
Ming          400
Felicia       300
Toby          222
Name: Cousins, dtype: int64
```

```
In [43]: 1 """ or all at once """
          2 s2 = pd.Series(['02334','19284','11111','96823','55555'],
          3                  name='ZipCodes',
          4                  index=['Providence','New Jersey','Delaware','Kansas City','Selma'])
          5 print(s2)
          6 print("")
          7 print("Where's Kansas City?")
          8 print(s2['Kansas City'])
```

```
Providence    02334
New Jersey    19284
Delaware       11111
Kansas City    96823
Selma          55555
Name: ZipCodes, dtype: object
```

```
Where's Kansas City?
96823
```



```
In [44]: 1 s3 = pd.Series([50000,23423,982874,77558,99382,33423,12345,
2               83728,44412,150000],
3               name='Income',
4               index=['Sacramento','Davis','Yolo',
5                     'Lakeport','Taho','Oakland',
6                     'Berkeley','Merced','Bakersfield','Stockton'])
7 print("_"*50,"\n",s3)
8 print("\nMedian: ", s3.median())
9 print("Mean: ",s3.mean())
10 print("STD: ",s3.std())
11 print("\nWho has the most and least?")
12 print("The least: $", s3.min(), " and the most $", s3.max())
```

---

Sacramento	50000
Davis	23423
Yolo	982874
Lakeport	77558
Taho	99382
Oakland	33423
Berkeley	12345
Merced	83728
Bakersfield	44412
Stockton	150000


Name: Income, dtype: int64

Median: 63779.0  
Mean: 155714.5  
STD: 293496.96320073836

Who has the most and least?  
The least: \$ 12345 and the most \$ 982874

```
In [45]: 1 import matplotlib.pyplot as plt
2
3 fig, axes = plt.subplots(1, 4, figsize=(12,3))
4 s3.plot(ax=axes[0], kind='line', title='line')
5 s3.plot(ax=axes[1], kind='bar', title='bar')
6 s3.plot(ax=axes[2], kind='box', title='box')
7 s3.plot(ax=axes[3], kind='pie', title='pie')
```

Out[45]: <AxesSubplot:title={'center':'pie'}, ylabel='Income'>



## Data Frames

```
In [46]: 1 import pandas as pd
2
3 # make some data
4 df = pd.DataFrame([[100000, "London"],
5                    [500000, "Paris"],
6                    [250000, "Rome"],
7                    [200000, "New_York"]])
```

```
In [47]: 1 # more useful to add column names that make sense:
2 df.index = ["UK","France","Italy","USA"]
3 df.columns = ['Good_Bakeries', 'Country']
4 """ Ouput some tests """
5 print(df)
6
7 print("\n","-"*40)
8 print("Just the bakeries and cities")
9 print(df.Good_Bakeries)
10
11 print("\n","-"*40)
12 print("What about USA? (using .loc)")
13 print( df.loc["USA"] )
14
15 print("\n","-"*40)
16 print("Can we compare 2 countries? (using .loc)")
17 print( df.loc[["UK", "USA"]])
18
19 print("\n","-"*40,"\nLearn about the data frame ... ")
20 print(df.info())
21 print("\nMay be silly but how about some stats? mean,std, median, min, max, etc.?" )
22 print(df.mean())
```

	Good_Bakeries	Country
UK	100000	London
France	500000	Paris
Italy	250000	Rome
USA	200000	New_York

---

Just the bakeries and cities

UK	100000
France	500000
Italy	250000
USA	200000

Name: Good\_Bakeries, dtype: int64

---

What about USA? (using .loc)

Good_Bakeries	200000
---------------	--------

```
Country          New_York
Name: USA, dtype: object
```

```
-----
Can we compare 2 countries? (using .loc)
```

```
    Good_Bakeries  Country
UK             100000   London
USA            200000  New_York
```

```
-----
Learn about the data frame ...
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 4 entries, UK to USA
```

```
Data columns (total 2 columns):
```

#	Column	Non-Null Count	Dtype
0	Good_Bakeries	4 non-null	int64
1	Country	4 non-null	object

```
dtypes: int64(1), object(1)
```

```
memory usage: 256.0+ bytes
```

```
None
```

```
May be silly but how about some stats? mean,std, median, min, max, etc.?
```

```
Good_Bakeries    262500.0
```

```
dtype: float64
```

---

## Reading some data for demo

Let's get some fake real data from a .csv file (cities.csv) ... to imagine we've gotten real files from work.

NB: in our source file, strings should be in double-quotes, not single ones.

The fake data look like this:

```
Rank, City, Country, Population, CensusDate
1, London, United Kingdom, "9,000,000", June 2014
2, Berlin, Deutschland, "4,000,000", June 2014
3, Paris, République française, "1,500,000", Aug 2014
4, Marseilles, République française, "2,500,000", Aug 2014
5, München, Deutschland, "2,200,100", April 2014
6, 香港, 中国, "7,577,231", Oct 2021
```

```
In [48]: 1 df_pop = pd.read_csv("files/cities.csv", delimiter=",", encoding="utf-8", header=0)
2 print("First some info about our objects from the file and then the head()")
3 df_pop.info()
4 print("")
5 print(df_pop.head())
```

First some info about our objects from the file and then the head()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 6 entries, 0 to 5

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	Rank	6 non-null	int64
1	City	6 non-null	object
2	Country	6 non-null	object
3	Population	6 non-null	object
4	CensusDate	6 non-null	object

dtypes: int64(1), object(4)

memory usage: 368.0+ bytes

	Rank	City	Country	Population	CensusDate
0	1	London	United Kingdom	9,000,000	June 2014
1	2	Berlin	Deutschland	4,000,000	June 2014
2	3	Paris	République française	1,500,000	Aug 2014
3	4	Marseilles	République française	2,500,000	Aug 2014
4	5	München	Deutschland	2,200,100	April 2014

```
In [49]: 1 """ note that df.head(n) is the same as df[:n] """
```

```
Out[49]: ' note that df.head(n) is the same as df[:n] '
```

```
In [50]: 1 print("\n","-"*40,"\nNot happy with this messy data. Let's clean it up a bit by removing the
2 df_pop["NumericPopulation"] = df_pop.Population.apply(lambda x: int(x.replace(",","")))
3
4 df_pop["Country"] = df_pop["Country"].apply(lambda x: x.strip())
5 # and confirm
6 print(df_pop.head())
```

-----  
 Not happy with this messy data. Let's clean it up a bit by removing the comma; use apply to convert strings to ints; make a new column

	Rank	City	Country	Population	CensusDate	\
0	1	London	United Kingdom	9,000,000	June 2014	
1	2	Berlin	Deutschland	4,000,000	June 2014	
2	3	Paris	République française	1,500,000	Aug 2014	
3	4	Marseilles	République française	2,500,000	Aug 2014	
4	5	München	Deutschland	2,200,100	April 2014	

	NumericPopulation
0	9000000
1	4000000
2	1500000
3	2500000
4	2200100

```
In [51]: 1 print("\n","-"*50,"Lets index the data by the city name.  Change index by \"set_index()\" and
2 print("Sort by country and the city ... so set sort_index(level=0) (meaning first index); rep
3 df_pop2 = df_pop.set_index("City")
4 df_pop2 = df_pop2.sort_index()
5 print("just city: ", df_pop2.head())
```

----- Lets index the data by the city name. Change index by "set\_index()" and sort.

Sort by country and the city ... so set sort\_index(level=0) (meaning first index); replace 0 with whatever int you want

just city:	Rank	Country	Population	CensusDate	\
City					
Berlin	2	Deutschland	4,000,000	June 2014	
London	1	United Kingdom	9,000,000	June 2014	
Marseilles	4	République française	2,500,000	Aug 2014	
München	5	Deutschland	2,200,100	April 2014	
Paris	3	République française	1,500,000	Aug 2014	

City	NumericPopulation
Berlin	4000000
London	9000000
Marseilles	2500000
München	2200100
Paris	1500000



```
In [52]: 1 print("now country and city")
2 df_pop3 = df_pop.set_index(["Country", "City"]).sort_index(level=0)
3 df_pop3.head()
```

now country and city

Out[52]:

		Rank	Population	CensusDate	NumericPopulation
Country	City				
Deutschland	Berlin	2	4,000,000	June 2014	4000000
	München	5	2,200,100	April 2014	2200100
République française	Marseilles	4	2,500,000	Aug 2014	2500000
	Paris	3	1,500,000	Aug 2014	1500000
United Kingdom	London	1	9,000,000	June 2014	9000000

```
In [53]: 1 print("\nGet some counts by Country: how many cities appear for each country?")
2 city_count = df_pop.Country.value_counts()
3 print(city_count.head())
4
5 print("\nDropping, grouping, reducing ... Drop a column, groupby another column ... and can a
6 df_pop5 = (df_pop.drop("Rank", axis=1)
7             .groupby("Country").sum()
8             .sort_values("NumericPopulation",
9                           ascending=False))
10
11 print(df_pop5.head())
```

Get some counts by Country: how many cities appear for each country?

```
République française    2
Deutschland             2
United Kingdom          1
中国                   1
Name: Country, dtype: int64
```

Dropping, grouping, reducing ... Drop a column, groupby another column ... and can apply a reduction (like sum, mean...)

```
                NumericPopulation
Country
United Kingdom          9000000
中国                   7577231
Deutschland             6200100
République française    4000000
```

```
1 <hr />
2 <p>The section above is offered as a way of thinking about your data – exploring them and asking questions of them, useful preparatory activities before further analysis or ingesting the wrong data!</p>
```

End of the notebook.

